

LocalMind Technical Guide

Version 1.2.0

Transform Audio into Intelligence

Professional-grade transcription with AI-powered quality analysis.

100% offline. Zero cost. Complete privacy.

Table of Contents

- [Introduction](#)
- [System Requirements](#)
- [Installation & First Launch](#)
- [Section A: Transcription \(Speech-to-Text\)](#)
- [What is Transcription?](#)
- [Whisper Models Explained](#)
- [Language Support](#)
- [Transcription Settings](#)
- [Section B: LLM Quality Analysis](#)
- [What is LLM Analysis?](#)
- [LLM Provider Options](#)
- [Local LLM Models](#)
- [Cloud LLM Providers](#)
- [Quality Scoring Parameters](#)
- [Export Options](#)
- [Settings Reference](#)
- [Troubleshooting](#)
- [Privacy & Security](#)

Introduction

LocalMind is a desktop application that performs two distinct AI tasks:

Task	Technology	Purpose
Transcription	OpenAI Whisper	Convert speech to text
Quality Analysis	Local/Cloud LLM	Score and analyze conversations

These are separate systems that work together but can be used independently.

System Requirements

Minimum Requirements

Component	Requirement
Operating System	macOS 12 (Monterey) or later
RAM	8 GB
Storage	10 GB free space
Processor	Intel or Apple Silicon

Recommended Requirements

Component	Requirement
Operating System	macOS 14 (Sonoma) or later
RAM	16 GB or more
Storage	20 GB free space
Processor	Apple M1/M2/M3 chip

First Run Downloads

On first launch, LocalMind downloads AI models:

Model Type	Size	When Downloaded
Whisper (transcription)	~1.5 GB	First transcription
Local LLM (analysis)	~4 GB	First quality analysis

Internet required only for initial model downloads.

Installation & First Launch

Step 1: Download

Download LocalMind-1.2.0-macOS.dmg from:

github.com/KaivalyaDeepTeam/LocalMind/releases

Step 2: Install

- Open the downloaded DMG file
- Drag LocalMind to your Applications folder

- Eject the DMG

Step 3: First Launch

Important: macOS may block the app because it's not from the App Store.

To open LocalMind:

- Right-click on LocalMind.app
- Select "Open" from the menu
- Click "Open" in the security dialog

This is a one-time step. After this, you can open normally.

Step 4: Model Download

On first use:

- Transcription models download when you process your first audio file
- LLM models download when you enable quality scoring

Progress is shown in the status bar. This may take 5-15 minutes depending on your internet speed.

Section A: Transcription (Speech-to-Text)

This section covers converting audio to text using OpenAI's Whisper technology.

What is Transcription?

Transcription converts spoken words in audio files into written text. LocalMind uses OpenAI Whisper, one of the most accurate speech recognition systems available.

How It Works

Audio File → Whisper AI → Written Transcript

(MP3) (Local) (Text)

Key Features

- 50+ languages supported
- Automatic language detection
- Speaker identification (diarization)
- Timestamps for each segment
- Works completely offline after model download

Supported Audio Formats

Format	Extension	Description
MP3	.mp3	Most common format
WAV	.wav	Uncompressed, high quality

M4A	.m4a	Apple/iTunes format
FLAC	.flac	Lossless compression
OGG	.ogg	Open source format
WebM	.webm	Web audio format

Maximum file size: 2 GB per file

Whisper Models Explained

Whisper comes in different sizes. Larger models are more accurate but slower.

Available Models

Model	Size	Accuracy	Speed	Best For
Large V3	1.5 GB	97-99%	Slow	Professional use, critical accuracy
Medium	750 MB	95-97%	Medium	Daily use, good balance
Small	250 MB	92-95%	Fast	Quick transcriptions
Base	150 MB	88-92%	Very Fast	Testing, drafts
Tiny	75 MB	80-88%	Fastest	Real-time, low accuracy OK

Model Recommendation

Your Priority	Recommended Model
Best accuracy	Large V3
Balanced performance	Medium
Fast processing	Small
Limited disk space	Base

How to Change Model

- Go to Settings (Cmd + ,)
- Select Transcription tab
- Choose model from dropdown
- Model downloads automatically if not present

Transcription Language Support

Automatic Detection

By default, LocalMind automatically detects the spoken language. This works well for:

- Single-language recordings
- Clear audio quality

Manual Selection

For best results, manually select the language when:

- Audio has background noise
- Multiple languages are mixed
- Accuracy is critical

Supported Languages (50+)

European: English, Spanish, French, German, Italian, Portuguese, Dutch, Polish, Russian, Ukrainian, Czech, Greek, Swedish, Danish, Norwegian, Finnish

Asian: Chinese (Mandarin), Japanese, Korean, Hindi, Bengali, Tamil, Telugu, Thai, Vietnamese, Indonesian, Malay

Middle Eastern: Arabic, Hebrew, Turkish, Persian, Urdu

And many more...

Transcription Settings

Access via Settings → Transcription

Basic Settings

Setting	Options	Description
Model	Tiny → Large V3	Accuracy vs speed tradeoff
Language	Auto / Specific	Audio language selection
GPU Acceleration	On / Off	Use Apple Silicon GPU

Advanced Settings

Setting	Default	Description
---------	---------	-------------

Chunk Length	30 sec	Audio segments for processing
Batch Size	16	Parallel chunks (RAM dependent)
Compute Type	float16	Precision (float32 for CPU)

Performance Tips

For faster processing:

- Enable GPU Acceleration (M1/M2/M3 Macs)
- Use Medium or Small model
- Increase Batch Size if you have 16GB+ RAM

For better accuracy:

- Use Large V3 model
- Select specific language (don't use auto-detect)
- Use high-quality audio source

Section B: LLM Quality Analysis

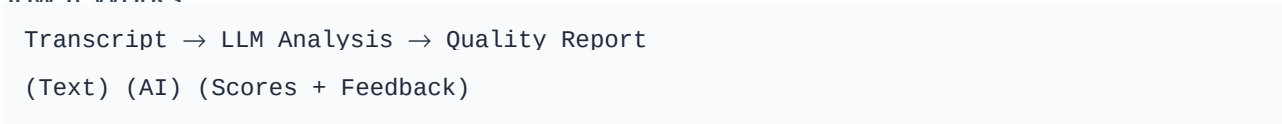
This section covers AI-powered conversation analysis using Large Language Models.

What is LLM Analysis?

LLM (Large Language Model) analysis reads your transcript and evaluates conversation quality. It provides:

- Overall score (0-100%)
- Parameter scores (customizable criteria)
- Strengths identified in the conversation
- Areas for improvement
- Detailed feedback for each parameter

How It Works



Key Difference from Transcription

Aspect	Transcription	LLM Analysis
Input	Audio file	Text transcript
Output	Written text	Scores & feedback

Technology	Whisper	LLM (Phi/Qwen/GPT)
Purpose	Convert speech	Evaluate quality
Required?	Yes	Optional

LLM Analysis requires transcription to be completed first.

LLM Provider Options

LocalMind supports three ways to run LLM analysis:

1. Local LLM (Recommended)

Runs entirely on your computer.

Pros	Cons
100% free	Slower than cloud
Complete privacy	Requires 8GB+ RAM
No internet needed	Large model download
No API limits	

2. OpenAI API

Uses OpenAI's GPT models via internet.

Pros	Cons
Very fast	Costs money (per use)
High quality	Requires internet
No local resources	Data sent to cloud

3. Anthropic API

Uses Anthropic's Claude models via internet.

Pros	Cons
Excellent reasoning	Costs money (per use)

Great for analysis	Requires internet
Detailed feedback	Data sent to cloud

How to Choose

Your Situation	Recommended Provider
Privacy is critical	Local LLM
No internet available	Local LLM
Need fastest results	OpenAI API
Budget available	OpenAI or Anthropic
Best analysis quality	Anthropic API

Local LLM Models

When using Local LLM, you can choose from several models:

Available Local Models

Model	Size	Speed	Quality	Best For
Phi-3.5 Mini	2.4 GB	Fast	Good	Default, balanced
Qwen 2.5 3B	2.0 GB	Very Fast	Good	Quick analysis
Qwen 2.5 7B	4.4 GB	Medium	Excellent	Professional use
Mistral 7B	4.1 GB	Medium	Excellent	Detailed feedback
Gemma 2 2B	1.6 GB	Fastest	Moderate	Speed priority

Model Recommendation

Priority	Recommended Model
Best quality	Qwen 2.5 7B or Mistral 7B
Balanced	Phi-3.5 Mini (default)

Fastest	Gemma 2 2B
Low RAM (8GB)	Qwen 2.5 3B

How to Change Local Model

- Go to Settings → LLM Provider
- Select Local LLM
- Click Model dropdown
- Choose your preferred model
- Model downloads automatically

Cloud LLM Providers

Setting Up OpenAI

- Get API key from platform.openai.com
- Go to Settings → LLM Provider
- Select OpenAI API
- Paste your API key
- Click Verify to test connection

Recommended model: GPT-4o Mini (best value)

Approximate cost: \$0.01-0.05 per audio file

Setting Up Anthropic

- Get API key from console.anthropic.com
- Go to Settings → LLM Provider
- Select Anthropic API
- Paste your API key
- Click Verify to test connection

Recommended model: Claude 3.5 Haiku (best value)

Approximate cost: \$0.01-0.03 per audio file

Quality Scoring Parameters

LocalMind evaluates conversations using customizable parameters.

Default Parameters

Parameter	Weight	What It Measures
Greeting & Introduction	1.0x	Professional opening
Active Listening	1.0x	Attention and engagement

Problem Identification	1.0x	Understanding the issue
Solution Provided	1.0x	Helpful resolution
Product Knowledge	1.0x	Accuracy of information
Communication Clarity	1.0x	Clear explanations
Empathy & Rapport	1.0x	Emotional connection
Call Control	1.0x	Managing conversation flow
Call Closing	1.0x	Professional ending
Script Compliance	1.0x	Following guidelines

Understanding Weights

Weights range from 0.1x to 3.0x:

Weight	Meaning	Impact on Score
0.1x - 0.5x	Low priority	Minor impact
1.0x	Standard	Normal impact
1.5x - 2.0x	High priority	Significant impact
2.5x - 3.0x	Critical	Major impact

Custom Scoring Profiles

Create profiles for different use cases:

Sales Calls:

- Product Knowledge: 2.5x
- Call Closing: 2.0x
- Communication Clarity: 2.0x

Support Calls:

- Empathy & Rapport: 2.5x
- Problem Identification: 2.0x
- Solution Provided: 2.5x

Compliance Audits:

- Script Compliance: 3.0x
- Greeting & Introduction: 2.0x

How to Customize

- Go to Edit → Scoring Parameters (Cmd + Shift + S)
- Adjust weight sliders
- Add/remove parameters as needed
- Save as new profile
- Select profile when processing

Export Options

LocalMind offers four export formats:

PDF Report (Recommended)

Keyboard: Cmd + Shift + P

Includes:

- Circular score gauge with overall percentage
- Bar chart visualization of all parameters
- Detailed scores table with ratings
- AI feedback (summary, strengths, improvements)
- Full transcript (optional)

Best for: Sharing with management, clients, or stakeholders

Markdown Report

Keyboard: Cmd + Shift + M

Includes:

- Formatted text with scores
- Strengths and improvements
- Easy to view in any text editor

Best for: Quick sharing, documentation, version control

JSON Export

Keyboard: Cmd + Shift + J

Includes:

- Complete structured data
- All scores and metadata
- Machine-readable format

Best for: Integration with other systems, batch processing

Text Transcript

Keyboard: Cmd + Shift + T

Includes:

- Plain text transcript
- Speaker labels
- Timestamps

Best for: Simple archival, text processing

Settings Reference

Transcription Tab

Setting	Description
Whisper Model	Speech recognition model size
Language	Audio language (Auto or specific)
GPU Acceleration	Use Apple Silicon for speed
Chunk Length	Audio segment size (advanced)
Batch Size	Parallel processing (advanced)

LLM Provider Tab

Setting	Description
Provider	Local LLM, OpenAI, or Anthropic
Model	Specific model to use
API Key	Required for cloud providers
Temperature	Creativity (0.0-1.0, lower = consistent)

Output Tab

Setting	Description
Output Directory	Where to save files
Auto-export JSON	Save JSON after processing
Auto-export PDF	Save PDF after processing
Include Transcript	Add transcript to PDF

Appearance Tab

Setting	Description
Theme	Light, Dark, or System
Language	UI language (8 available)
Colorblind Mode	Accessible color scheme

Troubleshooting

Transcription Issues

Problem: Transcription is very slow

Solutions:

- Enable GPU Acceleration in Settings
- Use a smaller Whisper model (Medium or Small)
- Close other applications to free RAM
- Process shorter audio files

Problem: Poor transcription accuracy

Solutions:

- Use Large V3 model
- Select specific language instead of Auto-detect
- Use higher quality audio source
- Enable audio preprocessing

Problem: Wrong language detected

Solution:

- Manually select the correct language in Settings

LLM Analysis Issues

Problem: Quality scoring not working

Solutions:

- Ensure LLM provider is configured in Settings
- Check if Local LLM model is downloaded
- Verify API key if using cloud provider
- Enable "Quality Scoring" checkbox before processing

Problem: Local LLM is slow

Solutions:

- Use a smaller model (Phi-3.5 or Gemma 2)
- Close other applications
- Consider using cloud provider for speed

Problem: API error with cloud provider

Solutions:

- Check API key is correct
- Verify account has credits/balance
- Check internet connection
- Try again later (rate limits)

General Issues

Problem: App won't open

Solution (macOS):

- Right-click LocalMind.app
- Select "Open"
- Click "Open" in dialog

Problem: Models not downloading

Solutions:

- Check internet connection
- Ensure sufficient disk space (10GB+)
- Check firewall/VPN settings
- Try restarting the app

Problem: Out of memory error

Solutions:

- Close other applications
- Use smaller models
- Restart the app
- Process shorter audio files

Privacy & Security

Data Handling

Mode	Audio Data	Transcript
Local LLM	Stays on device	Stays on device
OpenAI API	Stays on device	Sent to OpenAI
Anthropic API	Stays on device	Sent to Anthropic

Your audio files are NEVER uploaded to the cloud.

Only text transcripts are sent when using cloud LLM providers.

What LocalMind Collects

Nothing.

- No telemetry
- No analytics
- No crash reports
- No usage tracking
- No account required

Data Storage Locations

Data	Location
Transcripts	Your chosen output folder
Whisper models	~/.cache/whisper/
LLM models	~/.cache/localmind/models/
Settings	~/Library/Application Support/LocalMind/

Deleting All Data

To completely remove LocalMind data:

```
# Remove app
rm -rf /Applications/LocalMind.app

# Remove models
rm -rf ~/.cache/whisper/
rm -rf ~/.cache/localmind/

# Remove settings
rm -rf ~/Library/Application\ Support/LocalMind/
```

Keyboard Shortcuts

Action	Shortcut
Open Audio File	Cmd + O
Start Processing	Cmd + Return
Stop Processing	Escape
Export PDF	Cmd + Shift + P
Export Markdown	Cmd + Shift + M
Export JSON	Cmd + Shift + J

Export Transcript	Cmd + Shift + T
Scoring Parameters	Cmd + Shift + S
Settings	Cmd + ,
Show Shortcuts	Cmd + /
Quit	Cmd + Q

Getting Help

Resources

- Documentation: github.com/KaivalyaDeepTeam/LocalMind
- Issues: github.com/KaivalyaDeepTeam/LocalMind/issues
- Discussions: github.com/KaivalyaDeepTeam/LocalMind/discussions

Reporting Bugs

When reporting an issue, include:

- What you were trying to do
- What happened instead
- Your macOS version
- LocalMind version (1.2.0)
- Any error messages

About LocalMind

LocalMind is free, open-source software built to provide professional-grade audio analysis while respecting your privacy.

Our Principles

- Free Forever - No subscriptions, no hidden costs
- Privacy First - Your data stays on your device
- Open Source - Transparent, auditable code
- Offline Capable - Works without internet

Credits

- Whisper by OpenAI - Speech recognition
- Phi-3.5 by Microsoft - Local language model
- Llama.cpp - Efficient local inference
- PySide6 - Cross-platform UI
- ReportLab - PDF generation

Version: 1.2.0

Last Updated: January 2026

License: MIT (Free for any use)

© 2026 LocalMind Team. Made with care for everyone who values privacy.