

LocalMind Technical Guide

LocalMind

Contents

LocalMind	1
Technical Guide	1
Contents	1
Quick Start	2
Your First Transcription	2
Understanding Quality Scoring	3
Choosing the Right Model	3
Settings	4
Export & Share	6
Multilingual Support	6
Troubleshooting	7
Privacy & Security	7
Advanced Tips	8
System Requirements	9
Getting Help	9
About LocalMind	9

LocalMind

Technical Guide

Transform Audio into Intelligence

Professional-grade transcription with AI-powered quality analysis. 100% offline. Zero cost. Complete privacy.

Contents

- Quick Start
- Your First Transcription
- Understanding Quality Scoring
- Choosing the Right Model
- Settings
- Export & Share

Main Window

Figure 1: Main Window

- Troubleshooting
-

Quick Start

What You Need

- **macOS** 10.15 or later
- **4GB RAM** minimum (8GB recommended)
- **Audio file** in MP3, WAV, M4A, FLAC, OGG, or WEBM format

First Launch

1. Download LocalMind
2. Move to Applications folder
3. Double-click to open
4. Grant permissions if prompted

That's it. No account. No subscription. No internet required.

Your First Transcription

Step 1: Add Your Audio

Drag and drop your audio file into the window.

Supported formats: MP3 · WAV · M4A · FLAC · OGG · WEBM

File size: Up to 2GB per file

Step 2: Configure Processing

Choose your preferences:

Processing Mode: - **Offline** - Processes locally on your device - **Online** - Uses cloud AI (requires API keys)

Language: Auto-detect or select from 50+ languages

Model: Large V3 (Best Quality) - Recommended for first use

Audio Preprocessing: Enable noise reduction for clearer results

Step 3: Process

Click **Process** and watch the pipeline:

1. **Transcription** - Converting speech to text

Scoring Parameters

Figure 2: Scoring Parameters

Model Selection

Figure 3: Model Selection

2. **Merge Channels** - Combining audio streams
3. **Quality Audit** - AI-powered analysis
4. **Generate Report** - Creating comprehensive output

Processing time: 10-minute audio 5-7 minutes on average laptop

Understanding Quality Scoring

LocalMind doesn't just transcribe—it evaluates your conversations using advanced AI reasoning.

Default Parameters

Compliance (1.0x weight) - Greeting & Introduction - Active Listening - Problem Identification - Solution Provided - Product Knowledge - Communication Clarity - Empathy & Rapport - Call Control - Call Closing - Script Compliance

Customizing Scores

Adjust parameter weights from 0.1x to 3.0x:

- **Higher weight** = More important to overall score
- **Lower weight** = Less impact on final rating

Example: For sales calls, increase “Product Knowledge” to 2.5x

How Scoring Works

LocalMind uses **Chain of Thought (CoT) reasoning**:

1. Analyzes full transcript context
2. Identifies key moments and patterns
3. Evaluates against each parameter
4. Provides detailed explanations
5. Calculates weighted final score

Result: Understand not just *what* was said, but *how well* it was communicated.

Choosing the Right Model

Transcription Models

Qwen 2.5 (7B) - Best for auditing (**Recommended**)

LLM Provider Settings

Figure 4: LLM Provider Settings

- **Size:** 4GB
- **Speed:** Fast
- **Quality:** Excellent JSON output
- **Best for:** Quality analysis, professional use

Qwen 2.7B (6.4GB) - High quality audio

- **Size:** 6.4GB
- **Speed:** Moderate
- **Quality:** Very accurate for clear audio
- **Best for:** Structured transcription

Mixtral 7b-v3 (4.6GB) - Excellent output

- **Size:** 4.6GB
- **Speed:** Balanced
- **Quality:** Great reasoning ability
- **Best for:** Good all-around performance

Qwen 2.5 (3.2B) - Good balance

- **Size:** 3.2GB
- **Speed:** Faster
- **Quality:** Good for most use cases
- **Best for:** Smaller files, quick processing

Gemma 2 (2.6GB) - Very fast

- **Size:** 2.6GB
- **Speed:** Very fast
- **Quality:** Good for simple audio
- **Best for:** Fast turnaround needs

Whisper Models

Large V3 - Maximum accuracy (97-99%) **Medium** - Balanced performance (95-97%) **Base** - Speed priority (90-92%)

Settings

Access settings via **Settings** menu or **,** (Command-Comma)

Provider Dropdown

Figure 5: Provider Dropdown

Transcription Settings

Figure 6: Transcription Settings

LLM Provider

Choose your AI provider:

Local LLM (Free, Offline) - No internet required - Complete privacy - No API costs - Recommended for most users

OpenAI API - Requires API key - Pay per use - Cloud processing

Anthropic API - Requires API key - Advanced reasoning - Cloud processing

Transcription Settings

Model: Large V3 (Best Quality)

Language: Auto-detection Automatically identifies the spoken language

GPU Acceleration: Enable for 3-5x faster processing (if available)

Advanced Settings:

- **Chunk Length:** 30 seconds (default)
- **Batch Size:** 16 (adjust based on RAM)

Output Settings

Output Directory: Choose where to save results

Auto-export after processing: - Auto-export JSON - Auto-export PDF

PDF Report Settings: - Include full transcript - Include score breakdown

Appearance

UI Language: English · Español · · · · Français ·

Theme: Dark · Light · System

Accessibility: - Colorblind-friendly mode Uses blue/purple/orange colors for score gauges

Output Settings

Figure 7: Output Settings

Appearance Settings

Figure 8: Appearance Settings

Language Dropdown

Figure 9: Language Dropdown

Export & Share

Available Formats

JSON Machine-readable data with full transcript and scores

PDF Professional report with formatting and visualizations

TXT Plain text transcript only

Exporting

1. Complete processing
2. Click **Export** button
3. Choose format(s)
4. Select destination
5. Click **Save**

Files are named automatically: `filename_transcript_2026-01-18.pdf`

Multilingual Support

LocalMind speaks your language.

Supported UI Languages

- **English**
- **Español** (Spanish)
- (Japanese)
- (Arabic) - with RTL layout
- (Hindi)
- (Russian)
- **Français** (French)
- (Chinese Simplified)

Change Language

Settings → Appearance → UI Language

Changes apply immediately. No restart required.

Transcription Languages

LocalMind transcribes **50+ languages** including:

English · Spanish · French · German · Italian · Portuguese · Dutch · Russian · Arabic · Hindi · Japanese · Korean · Chinese · and many more

Troubleshooting

Processing Takes Too Long

Try this: - Use smaller Whisper model (Medium instead of Large) - Enable GPU acceleration in Settings - Close other applications to free RAM - Process shorter audio segments

Low Quality Scores

Remember: - Quality scoring requires LLM to be downloaded - First run downloads models (may take time) - Ensure “Enable Quality Scoring” is checked - Check that audio quality is good

Audio Not Loading

Check: - File format is supported (MP3, WAV, M4A, FLAC, OGG, WEBM) - File size is under 2GB - File is not corrupted - You have read permissions for the file

App Won’t Open

macOS Security: 1. Right-click LocalMind 2. Select “Open” 3. Click “Open” in security dialog 4. Grant permissions if requested

Models Not Downloading

Check: - You have internet connection (for first download) - Sufficient disk space (models are 2-7GB each) - Firewall allows HuggingFace connections - No VPN blocking downloads

Privacy & Security

What LocalMind Collects

Nothing.

- No telemetry
- No analytics
- No crash reports
- No usage statistics

Your audio never leaves your device in offline mode.

Data Storage

All data stored locally: - **Transcripts:** Your chosen output directory - **Models:** ~/.cache/localmind/
- **Settings:** ~/Library/Application Support/localmind/

Open Source

LocalMind is open source (MIT License).

Audit the code yourself: github.com/prepladder/localmind

Advanced Tips

Optimize Processing Speed

1. Use **GPU acceleration** if you have a Mac with M-series chip
2. **Choose appropriate model size** - Medium is sufficient for most needs
3. **Increase batch size** in advanced settings (if you have 16GB+ RAM)
4. **Process during off-hours** for background operation

Improve Transcription Accuracy

1. Use **highest quality audio** possible
2. **Enable audio preprocessing** for noisy recordings
3. **Select correct language** instead of auto-detect
4. Use **Large V3 model** for critical transcriptions

Batch Processing

Process multiple files efficiently:

1. Process first file with desired settings
2. Settings are remembered for next file
3. Enable auto-export to save time
4. Use same output directory for organized results

Custom Scoring Profiles

Create profiles for different use cases:

Sales Calls: - Product Knowledge: 2.5x - Communication Clarity: 2.0x - Call Closing: 2.0x

Support Calls: - Empathy & Rapport: 2.5x - Problem Identification: 2.0x - Solution Provided: 2.5x

Compliance Audits: - Script Compliance: 3.0x - Greeting & Introduction: 2.0x - Call Closing: 2.0x

System Requirements

Minimum

- macOS 10.15 Catalina or later
- 4GB RAM
- 10GB free disk space
- Intel or Apple Silicon processor

Recommended

- macOS 12 Monterey or later
 - 8GB RAM or more
 - 20GB free disk space
 - Apple M-series chip (for GPU acceleration)
-

Getting Help

Documentation

Full documentation: docs.localmind.ai

Community

- GitHub Issues: github.com/prepladder/localmind/issues
- Discussions: github.com/prepladder/localmind/discussions

Contact

- Email: support@localmind.ai
 - Website: localmind.ai
-

About LocalMind

LocalMind was built to give everyone access to professional-grade transcription and quality analysis without sacrificing privacy or paying monthly subscriptions.

Our Promise:

- Always free
- Always offline-capable
- Always open source
- Always privacy-focused

Version: 1.0.0 **Last Updated:** January 2026

Made with care for researchers, podcasters, journalists, call centers, legal professionals, and anyone who values their privacy.

© 2026 LocalMind. Released under MIT License.